



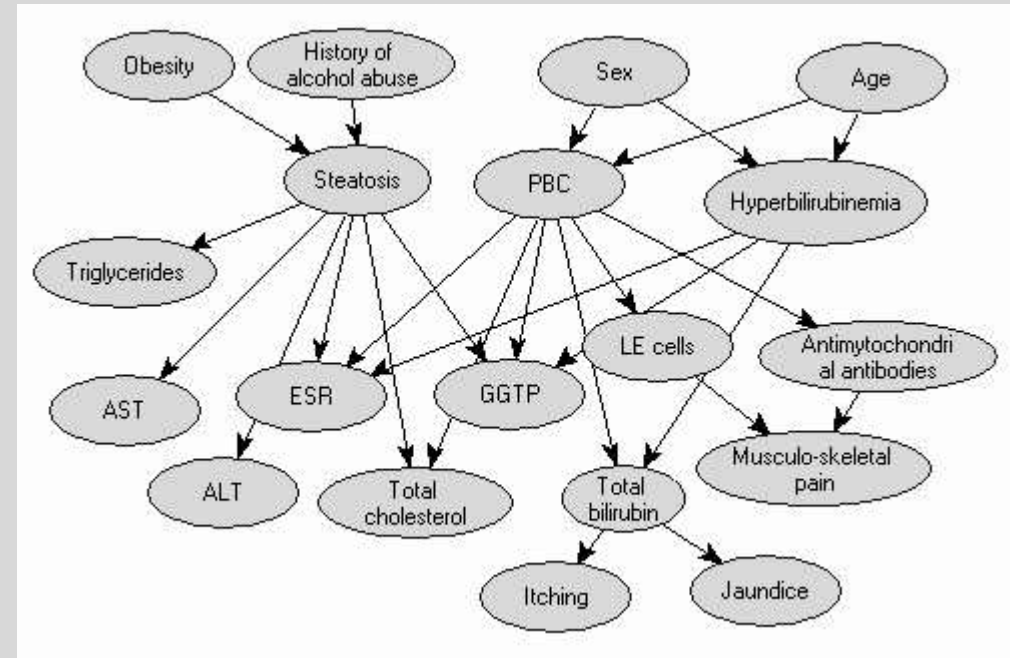
K NEAREST NEIGHBORS

Paul Speaker

Bayesian Wrap-up

- We have only scratched the surface of Bayesian Classification Models
- Steps in complexity
 - Naïve Bayes (no dependence)
 - LDA and QDA $\rightarrow P(X_1 | X_2)$
- Next steps: further dependencies
 - $P(X_1 | (X_3 | X_2))$ for example
 - **Bayesian Networks**
 - Many applications, particularly NLP

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$



K-Nearest Neighbors

- K nearest neighbors (KNN) is another machine learning algorithm used for classification.
- Basic idea: KNN works by identifying the K “nearest” data points to a given test data point and using them to determine the class or value of the test data point.
 - Has to use a train/test split
 - No model to build with train dataset
 - Simply make predictions on test dataset with k-nearest from training dataset
- Non-parametric—no explicit parameters for model
- Of all classification models, hardest to interpret probabilistically

“Nearest”

- “Nearest” is a more flexible concept than you might think at first
 - First requirement—“nearest” means only use continuous variables
 - Sometimes, binary categorical variables can be converted to 0/1 for purposes of KNN
- Problem: different X's have different scales
 - If one X has much larger spread (standard deviation) than others, it will dominate the distance calculation
 - Solutions: re-scaling X's
 - Normalized $X \rightarrow X' = (X - \mu)/\sigma$
 - (Note: subtracting mean does not really affect KNN, but it's often a standard part of scaling algorithms)
- Another potential issue—multiple ways of measuring distance
 - Standard—square root (sum of squares of differences) – like RMSE
 - “Taxicab”—sum of (absolute value of differences) – like MAE

KNN in R

- Command to do KNN is...`knn`
- Format:
 - `knn(train = <train dataset>, test = <test dataset>, cl = <class field in train>, k=<number>)`
- Scaling X's can be done with `scale()` command
 - before putting into `knn`,
 - `train_scale = scale(train)` `test_scale = scale(test)`
 - Or scale before split
- Other options:
 - `prob = <TRUE/FALSE>`
 - Default is FALSE
 - If TRUE, returns proportion of K nearest that are one class or other (sort of a probability)
 - `use.all` → handling of ties (for binary target, best to use odd K)
 - `l` = (deviations from majority rule, like a classification threshold)

Advantages and Limitations

- Advantages of KNN
 - Intuitive (I've never seen KNN rejected because results didn't make sense)
 - Works especially well for multi-valued (more later)
 - Non-parametric means highly nonlinear relationships can be captured
- Limitations of KNN
 - Suffers particularly hard from curse of dimensionality
 - Struggles more than most models with large numbers of inputs
 - Multivalued categorical inputs not really useful
 - Computationally expensive for larger datasets (search for nearest is hard)
 - Also struggles more than most with imbalanced datasets

Enhancements

- Most important enhancement to KNN is weighted KNN
- Idea is to weight closer point more heavily
 - Even if 4 out of 5 closest points are one category, it is possible to classify as second category if 5th point is close enough
- Calculate with weighted average
 - Weights are usually reciprocal of the distance
- Another option: “all within”
 - Rather than pick K and computing the closest, you pick a distance and pick all points within the fixed distances
 - Still can use distance weighting
 - My experience: often outperforms KNN
 - In R: `frNN` (does not weight by distance)
 - Python: `sklearn.neighbors.RadiusNeighborsClassifier` (can be weighted or not)